

SHUBH MEHTA

Open to Relocation | +1-812-803-4508 | shubh9mehta@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data analyst with an M.S. in Data Science who has worked every link of the data chain: cleaning messy CRM and warehouse data, writing the SQL and ETL behind it, building forecasting and classification models, and turning it all into Power BI and Tableau dashboards that teams actually open. Also ships AI systems, including a GPT-4o retrieval assistant at Indiana University that reached 90%+ semantic search precision. Happiest when raw, imperfect data turns into a decision someone can act on.

TECHNICAL SKILLS

Languages & Querying: Python (pandas, NumPy, scikit-learn), SQL (PostgreSQL, T-SQL, MySQL, BigQuery), window functions, CTEs, Excel (PivotTables, INDEX-MATCH, SUMIFS)

Data Engineering & Cloud: ETL pipeline design, data validation, star schema modeling, Mage-AI, Docker, CI/CD (GitHub Actions, CodePipeline), REST and GraphQL APIs, GCP (BigQuery, Cloud Storage, Cloud Scheduler), AWS (S3, Elastic Beanstalk, ECR), Azure

Analytics & Machine Learning: Regression and classification (XGBoost, Random Forest, Gradient Boosting), time-series forecasting, statistical testing, GridSearchCV tuning, model evaluation (R^2 , RMSE, ROC-AUC, recall)

BI & Visualization: Power BI (DAX), Tableau, Looker Studio, Streamlit, GA4, matplotlib, seaborn, Plotly

AI & LLMs: OpenAI GPT-4o API, Retrieval-Augmented Generation (RAG), SentenceTransformers embeddings, semantic search, prompt engineering, LangChain

EXPERIENCE

Macrotech Corporation | *Data Analyst* | Aug 2025 – Present | Irving, TX

- Write SQL with CTEs and window functions over order, shipment, and install tables to trace 400+ networking hardware SKUs (routers, gateways, switches, kiosks) from purchase order to on-site install across 900+ hospitality properties worldwide
- Automate CRM data extraction in Python (requests, pandas), pulling order and inventory records via API into analysis-ready datasets and cutting the weekly inventory-to-CRM reconciliation from a 2-hour manual Excel session to a scripted run that finishes in minutes
- Build Power BI dashboards with DAX measures and scheduled refresh tracking stock levels, installation progress, and usage trends by product and property, giving procurement and operations a live view instead of static spreadsheets
- Forecast replenishment demand in Excel from usage run-rates (SUMIFS, INDEX-MATCH, PivotTables), reconciling CRM counts against physical inventory and flagging low stock before installs slip

CruiseAlly | *Growth Data Analyst* | Jun 2025 – Aug 2025 | Remote

- Built progressive fallback logic and a similarity-based recommendation engine in a Django + GraphQL backend so no-match cruise searches relax nights, dates, and regions to surface alternatives, reducing zero-result searches 28%
- Implemented GraphQL resolvers for cabin availability, holds, pricing, and booking flows, with structured error handling around the third-party Revelex API to keep bookings stable when upstream calls failed
- Audited the sitemap, JSON-LD structured data, and Open Graph tags, fixed GA4 internal-traffic filtering so engagement metrics reflected real users, and delivered page-level recommendations in stakeholder-ready reports

Indiana University – Research Data Commons | *NLP Research Assistant* | Jan 2025 – May 2025 | Bloomington, IN

- Built a Retrieval-Augmented Generation (RAG) assistant with Flask, GPT-4o, and SentenceTransformer embeddings over 10,000+ research records from 7+ university departments, reaching 87% retrieval precision vs. a 34% keyword-search baseline
- Wrote the ingestion pipeline (pandas, openpyxl, python-docx) converting 15+ scattered Excel and Word sources (2.3M tokens) into one JSON knowledge base with source-level metadata, cutting manual data processing time 60%
- Tuned cosine similarity thresholds against a 200-query validation set and ran feedback rounds with librarians and faculty, cutting average lookup time from ~5 minutes to under 30 seconds while curbing hallucinated answers

Indiana University – Kelley School of Business | *Data Analyst – Research Assistant* | May 2024 – Aug 2024 | Bloomington, IN

- Linked 22,383 IPO records to 90,000+ blockholder ownership records that shared no clean identifier, combining Python fuzzy matching with Stata validation to reach a 94% match rate (21,000 records), ~40% faster via vectorized batching
- Ran logistic regression with 1,000-iteration bootstrap robustness checks on 8,200 firm-year observations, finding founder-led firms 17% less likely to lay off employees post-IPO (odds ratio 0.83, $p < 0.01$); built the visuals used in paper drafts

TeknoBiz Solutions Pvt Ltd. | *Data Analyst Intern* | May 2022 – Sep 2022 | India

- Consolidated 100,000+ sales records scattered across regional Excel and CSV files into a structured SQL Server database with T-SQL standardization and validation checks, cutting manual prep work ~30%
- Built Random Forest and Gradient Boosting forecasting models lifting sales-forecast R^2 from 0.65 to 0.78, and Tableau dashboards that moved monthly reporting from a 2–3 week cycle to hours, contributing to an estimated \$50K in annual savings

PROJECTS

End-to-End ML Pipeline with Cloud Deployment

- Engineered a modular pipeline (ingestion from CSV, MongoDB, and S3 → transformation → training → inference) comparing 6 algorithms with GridSearchCV and 3-fold cross-validation; best model reached $R^2 = 0.87$, with saved preprocessing artifacts keeping training and live inference consistent
- Containerized with Docker and deployed twice over: AWS Elastic Beanstalk via CodePipeline + ECR, and Azure Web App via GitHub Actions + ACR, with a Flask front end serving real-time predictions

NYC Taxi Revenue & Operations Analytics (GCP)

- Designed a fact-dimension model and Mage-AI ETL pipeline loading 150,000 sampled NYC TLC trip records through Cloud Storage into BigQuery, with Cloud Scheduler automating refreshes end to end
- Analyzed \$4M+ in trip revenue through Looker Studio dashboards, surfacing 3–7 AM airport-driven peak hours and an 82.4% credit-card payment share, informing dispatch and driver-allocation decisions

EDUCATION

Indiana University Bloomington — M.S., Data Science | Aug 2023 – May 2025

SRM Institute of Science & Technology — B.Tech, Computer Science & Engineering | Jul 2019 – May 2023